

[Open in app ↗](#)

# Medium



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



# AI Agent Architecture: Mapping Domain, Agent, and Orchestration to Clean Architecture

5 min read · Aug 28, 2025




Naoyuki Sakai

Following ▾

Listen

Share

More



Generated by Naoyuki Sakai — AI Agent Architecture (Domain / Agent / Orchestration / Infrastructure)

## Introduction

AI agents aren't magic black boxes — they fit into Clean Architecture.

AI agents are often portrayed as a *black box*.

An LLM “reads” your request, “decides” which tools to call, and somehow produces the right answer.

This story is attractive but misleading. In real-world enterprise systems, software must be **auditable, testable, and maintainable**. You cannot rely on a monolithic prompt to encode business logic, nor can you let a stochastic model silently control mission-critical workflows.

To bring AI agents into the same engineering discipline as other systems, we need a clear architectural perspective.

In this article, I propose mapping AI agents onto **Clean Architecture / DDD principles**, showing how **Domain, Agent, and Orchestration** fit into the concentric layers of software design.

*Example upfront:* Think of a logistics company asking an AI agent, “*Generate tomorrow’s delivery plan.*” The answer should not come from a black box but from a transparent pipeline that interprets intent, gathers data, and validates against strict rules.

## Three Roles Defined

When building AI agent systems, three distinct responsibilities emerge:

### 1. Domain (Correctness)

- The core business truth: legal rules, business constraints, KPI formulas.
- Must be implemented in deterministic code and tested explicitly.
- Example: “Delivery trucks must not exceed 8 hours of operation per day.”

### 2. Agent (Tactical Selection)

- Interprets the user’s intent and chooses which use case or pipeline to execute.
- Think of it as the “router” that decides *what to do next*.
- Example: For a request “*Generate tomorrow’s delivery plan*”, the agent decides to call `GeneratePlan` and then `EvaluatePlan`.

### 3. Orchestration (Information Flow)


- Manages how data and tools are connected.

- Fetches relevant context, runs retrieval-augmented generation (RAG), merges results, and builds prompts.
- Example: Query demand forecast DB → retrieve past exceptions from VectorDB → fetch regulations from FAQ → fuse into context for the LLM.

👉 Put simply:

- **Domain = correctness**
- **Agent = tactical selection**
- **Orchestration = information flow**

### Visualizing the Separation



### Mapping to Clean Architecture

Clean Architecture, popularized by Robert C. Martin (Uncle Bob), organizes software into concentric layers:

- **Domain** (Entities, Use Cases) at the core
- **Application** controlling flow and orchestration
- **Infrastructure = Interface Adapters + Frameworks & Drivers**
  - *Interface Adapters*: bridges that connect the application to external systems (DB/VectorDB adapters, API adapters, UI adapters, retrievers, rerankers, prompt builders, etc.)
  - *Frameworks & Drivers*: the actual external technologies (LLM APIs, databases, vector stores, web frameworks, schedulers, monitoring systems, etc.)

Two timeless principles apply:


1. Inner layers must not depend on outer layers.

2. The core domain should remain stable and testable, even if technologies change outside.

## How AI Agents Fit

- **Domain Layer:** Immutable rules (constraints, KPI formulas, validators). Always implemented in code, never in prompts.
- **Application Layer:** Where AI agents live.
  - *Agent:* interprets intent and selects use cases.
  - *Orchestration:* arranges flows (retrieval, API calls, prompts).
- **Infrastructure Layer (Interface Adapters + Frameworks & Drivers):**
  - *Interface Adapters:* DB adapters, VectorDB adapters, API adapters, UI adapters, retriever/reranker adapters, prompt builders.
  - *Frameworks & Drivers:* LLM APIs, DB/VectorDB engines, web/UI frameworks, schedulers, monitoring/logging systems.

## Diagram (Mermaid)



- This keeps the original structure but clarifies:

- **Infrastructure = Interface Adapters + Frameworks & Drivers**
- **Application = Agents + Orchestration**

- Domain = Immutable business rules

## Where ChatGPT and RAG Fit


LLMs like ChatGPT are often mistaken as “the whole agent.” In fact:


### ChatGPT / LLMs

- Not the Domain.
- Instead:
  - Agent-LLM interprets intent and decides what to do.
  - Orchestration-LLM fuses data and generates natural-language output.
- ⚠️ LLMs are powerful but stochastic. They excel at interpretation and expression, but should not hold business-critical truth.

### RAG

- Not a new layer, but an orchestration pattern:
  - Retriever → Reranker → Fusion → Prompt → LLM





## Design Guidelines (Do)

To keep AI agents maintainable, auditable, and trustworthy:

### Must-have

1. **Externalize prompts** — version-controlled templates.
  2. **Separate Planner and Executor** — LLM proposes, code executes.
  3. **Use JSON Schema contracts** — validate tool I/O.
  4. **Ensure observability** — log input, output, rationale, model ID.
- Example log entry:

```
{
  "input": "Generate plan for Tokai",
  "output": {"planId": 123, "score": 0.87},
  "rationale": "Best load balance",
  "model": "gpt-4",
  "timestamp": "2025-08-25T10:00Z"
}
```

### Recommended

5. **Separate Router and Orchestrator** — avoid mixing intent routing with flow design.
6. **Declare orchestration pipelines** — YAML/DAG for RAG flows.

```
pipeline:
  - stage: retriever
```

- stage: reranker
- stage: fusion
- stage: prompt\_construction
- stage: llm\_answer

## 7. Keep the UI thin — presentation only.

### Anti-patterns (Don't)

Avoid these traps:

- **God Prompt** — one giant prompt that mixes everything.
- **Domain in Prompt** — hiding business rules in LLM instructions.
- **Tool Zoo** — unstructured tools with no capability map.
- **Fat UI/Application** — logic and prompts leaking into UI.
- **Implicit State** — hidden history, no explicit session store.


*(Common case: embedding prompts directly in UI code → impossible to test or reuse later.  
Seen frequently in early prototypes. Works short-term, breaks long-term.)*

### Example: Logistics Dispatch

A dispatch planner asks: “Generate tomorrow’s delivery plan for the Tokai region.”

- **Agent:** Interprets intent, decides GeneratePlan → EvaluatePlan.
- **Orchestration:** Retrieves demand forecast (DB), past exceptions (VectorDB), regulations (FAQ); fuses context and asks LLM to draft plan.
- **Domain:** Enforces rules (driver hours, load capacity, delivery windows). Invalid plans are rejected; valid ones committed via CommitPlan.

### Visualizing the Orchestration Flow



👉 Outcome: a transparent, auditable workflow — not a black box.

👉 Business benefit: improves efficiency *and* guarantees compliance with operational rules.

👉 Non-functional benefit: easier auditability and component reuse.

## Conclusion

AI agents do not require a brand-new architectural paradigm.

They can be **cleanly mapped** into existing **Clean Architecture principles**:

- **Domain = correctness**
- **Agent = tactical selection**
- **Orchestration = information flow**

With this separation, AI agents become **auditable, testable, and maintainable components** in enterprise systems — not magical black boxes, but reliable architecture citizens.

*Coming next in Part 2: How to handle multi-stage pipelines (retrieval → rerank → generation → validation) and fit them into this architecture.*